


EntroQuant vs prior methods — KV cache compression of R1-Distill

Information-theoretic bound: how much can be saved?



At $b = 1$ bit/coord: how wrong is $\langle q, k \rangle$ after compression?

